

D209 – Data Mining (Task 1)

Morrell J. Parrish

Western Governors University

## Table of Contents

A. RESEARCH QUESTION .....	3
A2. ANALYSIS GOAL .....	3
B. CHOSEN TECHNIQUE .....	3
B2. ASSUMPTION.....	3
B3. PACKAGES/LIBRARIES USE JUSTIFICATION .....	4
C. DATA PREPARATION DESCRIPTION .....	4
C2. VARIABLE IDENTIFICATION AND CLASSIFICATION .....	5
C3. DATA PREPARATION STEPS .....	5
C4. CLEANED DATA SET .....	7
D. ANALYSIS .....	7
D2. ANALYTICAL TECHNIQUE DESCRIPTION .....	7
D3. CLASSIFICATION ANALYSIS CODE.....	11
E. SUMMARY.....	13
E2. CLASSIFICATION ANALYSIS RESULTS AND IMPLICATIONS.....	14
E3. ANALYSIS LIMITATION .....	14
E4. RECOMMENDED COURSE OF ACTION .....	14
G. PANOPTO VIDEO RECORDING.....	15
REFERENCES.....	16

## D209 – Data Mining (Task 1)

**A. Research Question**

During this course of research, we will determine which customers are at a higher risk of churn and which features (variables) can be an indicator for churn.

**A2. Analysis Goal**

The objective of this analysis is to use data classification to determine which customers, new or existing, are at a higher risk for churn based on the similarities to other customers with similar features that may have or have not churned in the past. “The churn rate, also known as the rate of attrition or customer churn; is the frequency in which consumers discontinue doing business with a company. It is commonly represented as the percentage of service subscribers who cancel their memberships within a specified time frame” (Frankenfield, 2022).

**B. Chosen Technique**

The **k-nearest neighbors** classification method will be used for this analysis. **K-nearest neighbors** works on the principle that every data point that is close to another is in the same class. In other words, it uses similarity to classify a new data point (Dwivedi, *nd*)

The expected outcome of this classification analysis should show the targeted variable (churn) and the relationship between the nearing **k-neighbors**.

**B2. Assumption**

The primary assumption of a **KNN** model is that data points/instances that are close to each other are highly similar, whereas data points that are far away from another group are dissimilar to those data points. The distance between two points on a graph is used by a **KNN** model to calculate similarity. The farther apart the points are, the less similar they are. There are

several methods for calculating distance between points, but the most commonly used distance metric is Euclidean distance (the distance between two points in a straight line) (Nelson, 2020).

### **B3. Packages/Libraries Use Justification**

The following packages/libraries will be used for this analysis:

- Pandas
  - used to read and manipulate data via series (one-dimensional structure) or dataframes (multi-dimensional data structure)
- NumPy
  - used to perform mathematical computations
- Matplotlib
  - used to create visualization (plotting and graphing)
- Seaborn
  - used to create visualization (plotting and graphing)
- Scikit-learn
  - used to perform scientific computations
  - used to split our data into training and test sets
  - used for predicting and classification analysis

### **C. Data Preparation Description**

To use the churn dataset in our analysis we will first need to prepare the data.

The following steps were taken to prepare the dataset for analysis:

- download the churn dataset
- determine which variables will be used in the analysis
- import the dataset into *PyCharm*

- remove independent variables, demographics, and personal identification variables not being used in the analysis
  - caseorder, customer\_id, interaction, UID, city, state, county, zip, lat, lng, population, timezone, job, email, contacts
- determine if any outliers exist and remove them

## C2. Variable Identification and Classification

The **continuous variables** (16) that will be used in this analysis will include age, children, income, outage\_sec\_perweek, yearly\_equip\_failure, tenure, monthlycharge, bandwidth\_GB\_Year, item1 (timelyresponse), item2 (fixes), item3 (replacements), item4 (reliability), item5 (options), item6 (respectfulness), item7 (courteous), and item8 (listening).

The **categorical variables** (19) that will be used in this analysis will include area, marital, gender, churn, techie, contract, portmodem, tablet, internetservice, phone, multiple, onlinesecurity, onlinebackup, deviceprotection, techsupport, streamingtv, streamingmovies, paperlessbilling, paymentmethod.

## C3. Data Preparation Steps

To use the churn dataset in our analysis we will first need to prepare the data:

- import the dataset into *Python (PyCharm)*
- view the dataframe's description, structure, and data types
- view summary statistics
- evaluate the dataset, remove null or missing values
- remove any outliers
- remove demographics, and personal identification

- caseorder, customer\_id, interaction, UID, city, state, county, zip, lat, lng, population, area, timezone, job, email, contacts
- convert binomial variables (yes/no to 1 and 0) to numerical variables

**The below code was used to prepare our data:**

```
# Standard data science imports
import NumPy as np
import pandas as pd
from pandas import Series, DataFrame

# Visualization libraries
import seaborn as sns
import matplotlib.pyplot as plt

# Scikit-learn
import sklearn
from sklearn import datasets
from sklearn import preprocessing
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.metrics import classification_report

# Ignore Warning Code
import warnings
warnings.filterwarnings('ignore')

# Load data set into Pandas dataframe
df = pd.read_csv('churn_clean.csv')

# Remove less meaningful demographic variables
df = df.drop(columns=['CaseOrder', 'Customer_id', 'Interaction', 'UID',
'City', 'State', 'County', 'Zip', 'Lat', 'Lng', 'Population', 'TimeZone',
'Email', 'Contacts', 'Job'])

# Display Churn dataframe
print(df)

# Rename last 8 columns
df.rename(columns={'Item1': 'TimelyResponse', 'Item2': 'Fixes', 'Item3':
'Replacements', 'Item4': 'Reliability', 'Item5': 'Options', 'Item6':
'Respectfulness', 'Item7': 'Courteous', 'Item8': 'Listening'}, inplace=True)

# Get column info
print(df.info())

# Describe Churn dataset
print(df.describe())

# Convert binary variables (yes/no, female/male) to 0 or 1
df['DmyGender'] = [1 if v == 'Male' else 0 for v in df['Gender']]
df['DmyChurn'] = [1 if v == 'Yes' else 0 for v in df['Churn']]
df['DmyTechie'] = [1 if v == 'Yes' else 0 for v in df['Techie']]
```

```

df['DmyContract'] = [1 if v == 'Two Year' else 0 for v in df['Contract']]
df['DmyPort_modem'] = [1 if v == 'Yes' else 0 for v in df['Port_modem']]
df['DmyTablet'] = [1 if v == 'Yes' else 0 for v in df['Tablet']]
df['DmyInternetService'] = [1 if v == 'Fiber Optic' else 0 for v in df['InternetService']]
df['DmyPhone'] = [1 if v == 'Yes' else 0 for v in df['Phone']]
df['DmyMultiple'] = [1 if v == 'Yes' else 0 for v in df['Multiple']]
df['DmyOnlineSecurity'] = [1 if v == 'Yes' else 0 for v in df['OnlineSecurity']]
df['DmyOnlineBackup'] = [1 if v == 'Yes' else 0 for v in df['OnlineBackup']]
df['DmyDeviceProtection'] = [1 if v == 'Yes' else 0 for v in df['DeviceProtection']]
df['DmyTechSupport'] = [1 if v == 'Yes' else 0 for v in df['TechSupport']]
df['DmyStreamingTV'] = [1 if v == 'Yes' else 0 for v in df['StreamingTV']]
df['DmyStreamingMovies'] = [1 if v == 'Yes' else 0 for v in df['StreamingMovies']]
df['DmyPaperlessBilling'] = [1 if v == 'Yes' else 0 for v in df['PaperlessBilling']]

# Drop original categories
df2 = df.drop(columns=['Gender', 'Churn', 'Techie', 'Contract', 'Port_modem',
'Tablet', 'InternetService', 'Phone', 'Multiple', 'OnlineSecurity',
'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV',
'StreamingMovies', 'PaperlessBilling'])

print(df2.describe())

```

#### C4. Cleaned Data Set

The prepared dataset used for this analysis has been uploaded with the assessment file.

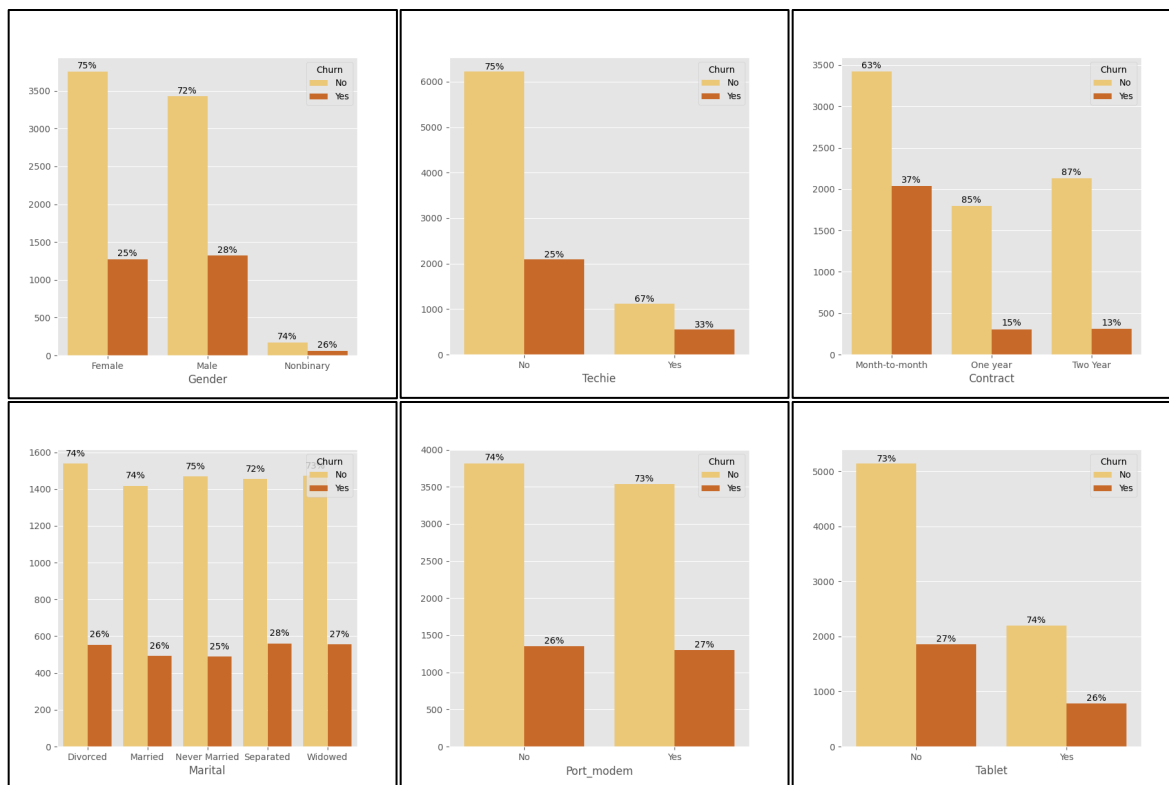
#### D. Analysis

The training and test datasets used for this analysis have been uploaded with the assessment file.

#### D2. Analytical Technique Description

Our analytical technique includes the following steps: (1) read in or load the data using Pandas' *read()* function – in this case it will be our cleaned churn data set, (2) check to make sure the data was read in properly using Pandas' *head()* function, (3) verify the shape of the data using Pandas' '*shape*' function, (4) perform exploratory analysis by creating a series of visual (boxplots, , scatterplots, correlation matrix, and heatmaps etc....), (5) identify and remove any outliers, (6) split up the dataset into inputs (X) and our target variable (y) using Pandas' *drop()* function – this allows you to drop the target variable from the dataframe and store it in the variable 'X', (7) verify the data again using Pandas' *head()* function, (8) split the dataset into training and test sets using Scikit-learn's function '*train\_test\_split*', (9) create/build the initial model using Scikit-learn's function '*KNeighborsClassifier*', (10) train model - this can be done

using the KNN's *fit()*, (11) once trained we can now test the model using KNN's *predict()* function, (12) verify model's accuracy on the test data using KNN's *score* function, (13) perform **k-Fold Cross Validation** – splits the data set into 'k', one group is used as the test set and the others are used as training sets – recording the accuracy score in an array for each test, (14) find the average of k-Fold Validation tests by using Numpy's *mean()* function, passing in the 'cv\_score' from the validation test, (15) perform hypertuning parameters using Scikit-learn's function *GridSearchCV()*, (16) check/verify after training which value for 'n\_neighbor' did the best by calling '*best\_params\_*' on the model, (17) after finding the optimal value, I used the '*best\_score\_*' function to check the accuracy of the model.





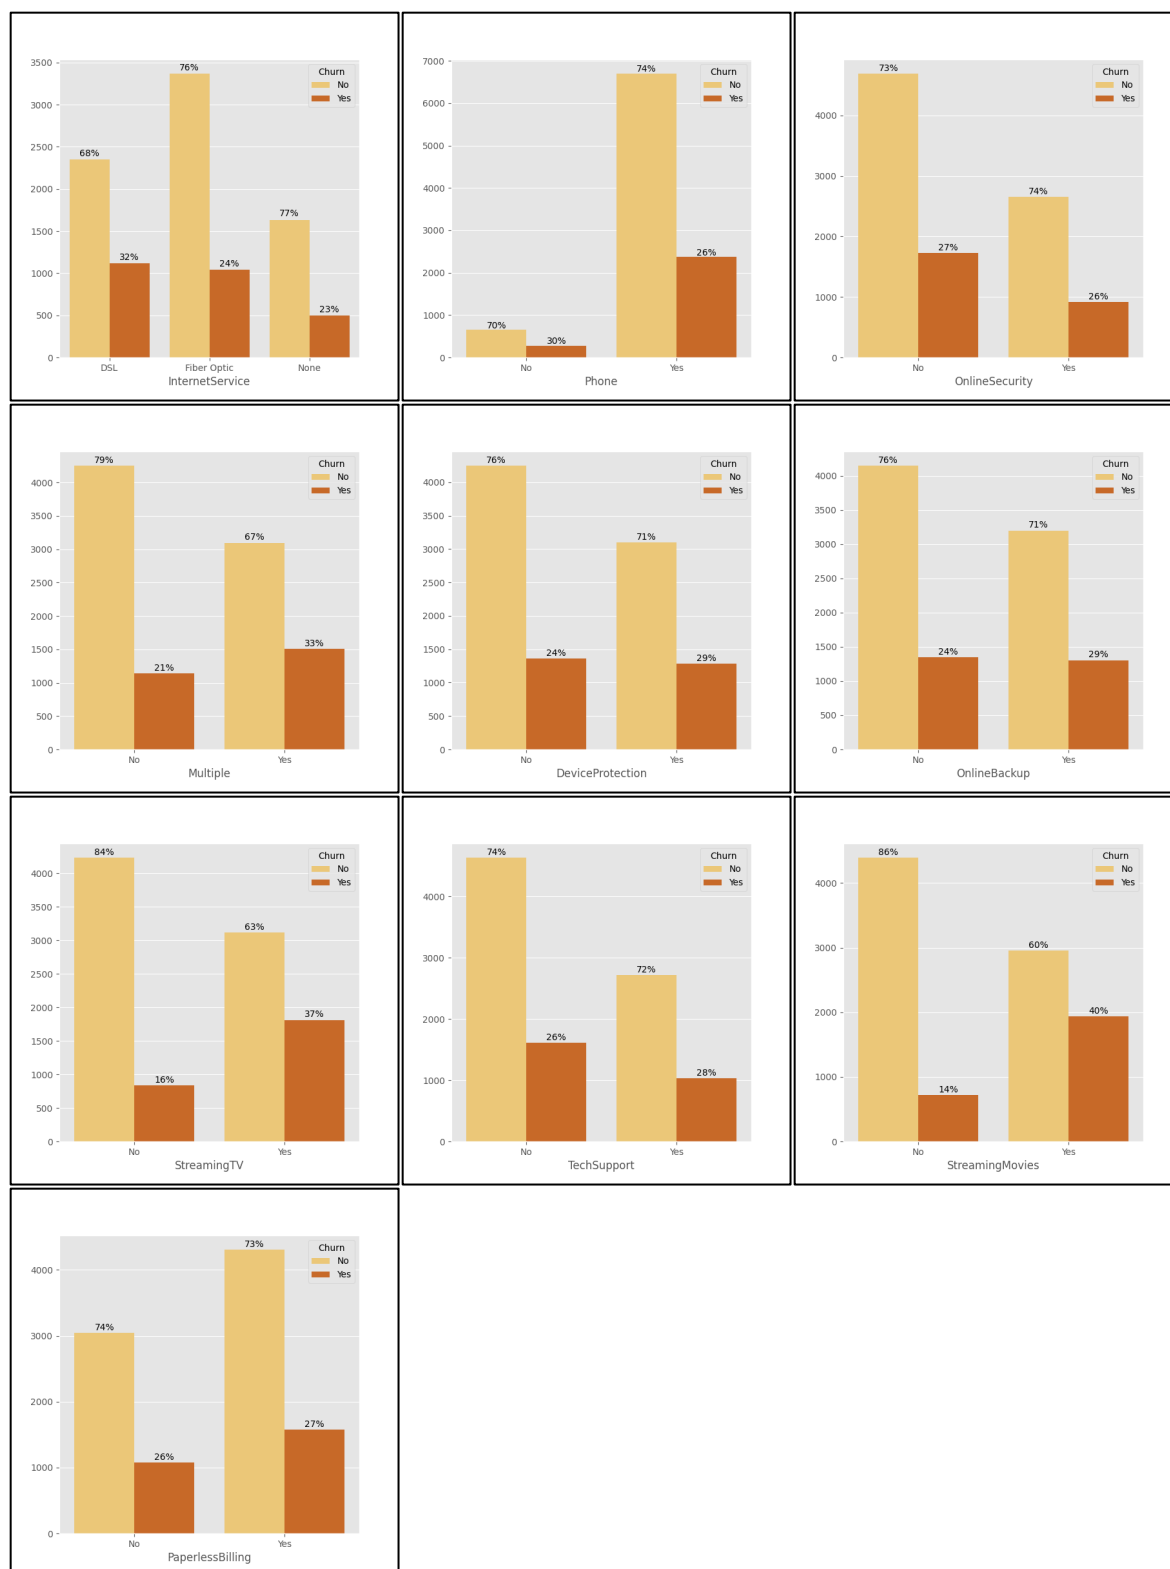


Figure 1: Churn and Categorical Variables

```

94 # Set predictor features & target variable
95 X = df4.drop('DmyChurn', axis=1).values
96 y = df4['DmyChurn'].values

103 # Create training and test sets
104 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20, random_state=SEED)
105
106 print("\nX_train:\n")
107 print(X_train.shape)
108
109 print("\nX_test:\n")
110 print(X_test.shape)
111
112 np.savetxt('X_train.txt', X_train)
113 np.savetxt('X_test.txt', X_test)
114
115 # Instantiate KNN model
116 knn = KNeighborsClassifier(n_neighbors=7)
117
118 # Fit data to KNN model
119 knn.fit(X_train, y_train)
120
121 # Predict outcomes from test set
122 y_pred = knn.predict(X_test)
123
124 # Print initial accuracy score of KNN model
125 print('Initial accuracy score KNN model: ', accuracy_score(y_test, y_pred))
126
127 # Compute classification metrics
128 print(classification_report(y_test, y_pred))

```

X\_train:

(8000, 31)

X\_test:

(2000, 31)

Initial accuracy score KNN model: 0.7085

	precision	recall	f1-score	support
0	0.78	0.83	0.80	1442
1	0.47	0.38	0.42	558
accuracy			0.71	2000
macro avg	0.63	0.61	0.61	2000
weighted avg	0.69	0.71	0.70	2000

Figure 2: Output and Intermediate Calculations




Figure 3: Confusion Matrix

According to the classification report our intermediate calculations yielded that 47% of the customers predicted to churn did so. The model also only correctly predicted this outcome for 38% of those customers; the model's accuracy score is about 71%.

```

159 # Set steps for pipeline object
160 steps = [('scaler', StandardScaler()),
161          ('knn', KNeighborsClassifier())]
162
163 # Instantiate pipeline
164 pipeline = Pipeline(steps)
165
166 # Split dataframe
167 X_train_scaled, X_test_scaled, y_train_scaled, y_test_scaled = train_test_split(X, y, test_size=0.2,
168                                                                                   random_state=SEED)
169
170 # Scale dataframe with pipeline object
171 knn_scaled = pipeline.fit(X_train_scaled, y_train_scaled)
172
173 # Predict from scaled dataframe
174 y_pred_scaled = pipeline.predict(X_test_scaled)
175
176 # Print new accuracy score of scaled KNN model
177 print('New accuracy score of scaled KNN model: {:.3f}'.format(accuracy_score(y_test_scaled,
178                                                                                   y_pred_scaled)))
179
180 # Compute classification metrics after scaling
181 print(classification_report(y_test_scaled, y_pred_scaled))

```

New accuracy score of scaled KNN model: 0.818					
	precision	recall	f1-score	support	
0	0.85	0.91	0.88	1442	
1	0.71	0.58	0.64	558	
accuracy			0.82	2000	
macro avg	0.78	0.74	0.76	2000	
weighted avg	0.81	0.82	0.81	2000	

I performed some hypertuning parameters on my initial model using Skicit-learn's **GridSearchCV()** function; this allowed our model to be trained (multiple times) on a range of specified parameters; after hypertuning our model our retrained scaled model calculations yielded that 71% of the customers predicted to churn did so. The model also only correctly predicted this outcome for 58% of those customers; the model's accuracy score is about 82%; the models accuracy increased by 11%.

### D3. Classification Analysis Code

```

# Set steps for pipeline object
steps = [('scaler', StandardScaler()),

```

```

('knn', KNeighborsClassifier())]

# Instantiate pipeline
pipeline = Pipeline(steps)

# Split dataframe
X_train_scaled, X_test_scaled, y_train_scaled, y_test_scaled =
train_test_split(X, y, test_size=0.2, random_state=SEED)

# Scale dataframe with pipeline object
knn_scaled = pipeline.fit(X_train_scaled, y_train_scaled)

# Predict from scaled dataframe
y_pred_scaled = pipeline.predict(X_test_scaled)

# Print new accuracy score of scaled KNN model
print('New accuracy score of scaled KNN model:
{:0.3f}'.format(accuracy_score(y_test_scaled, y_pred_scaled)))

# Compute classification metrics after scaling
print(classification_report(y_test_scaled, y_pred_scaled))

# Create confusion matrix & generate results
cf_matrix = confusion_matrix(y_test, y_pred)
print(cf_matrix)

# Generate confusion matrix visual
group_names = ['True Neg', 'False Pos', 'False Neg', 'True Pos']
group_counts = ["{0:0.0f}".format(value) for value in cf_matrix.flatten()]
group_percentages = ["{0:.2%}".format(value) for value in cf_matrix.flatten()
/ np.sum(cf_matrix)]
labels = [f"{v1}\n{v2}\n{v3}" for v1, v2, v3 in zip(group_names,
group_counts, group_percentages)]
labels = np.asarray(labels).reshape(2, 2)

sns.heatmap(cf_matrix, annot=labels, fmt='', cmap='Oranges')
plt.show()

# Set up parameters grid
param_grid = {'n_neighbors': np.arange(1, 50)}

# Prepare KNN for cross validation
knn = KNeighborsClassifier()

# Instantiate GridSearch cross validation
knn_cv = GridSearchCV(knn, param_grid, cv=5)

# Fit model to
knn_cv.fit(X_train, y_train)

# Print best parameters
print('Best parameters for this KNN model: {}'.format(knn_cv.best_params_))

# Generate model best score
print('Best score for this KNN model: {:.3f}'.format(knn_cv.best_score_))

# ROC AUC metrics for explaining the area under the curve

```

```

# Fit it to the data
knn_cv.fit(X, y)

# Compute predicted probabilities: y_pred_prob
y_pred_prob = knn_cv.predict_proba(X_test)[:, 1]

# Compute and print AUC score
print("The Area under curve (AUC) on validation dataset is:
{:.4f}".format(roc_auc_score(y_test, y_pred_prob)))

# Compute cross-validated AUC scores: cv_auc
cv_auc = cross_val_score(knn_cv, X, y, cv=5, scoring='roc_auc')

# Print list of AUC scores
print("AUC scores computed using 5-fold cross-validation: {}".format(cv_auc))

```

## E. Summary

```

200 # Prepare KNN for cross validation
201 knn = KNeighborsClassifier()
202
203 # Instantiate GridSearch cross validation
204 knn_cv = GridSearchCV(knn, param_grid, cv=5)
205
206 # Fit model to
207 knn_cv.fit(X_train, y_train)
208
209 # Print best parameters
210
211 print('Best parameters for this KNN model: {}'.format(knn_cv.best_params_))
212
213 # Generate model best score
214 print('Best score for this KNN model: {:.3f}'.format(knn_cv.best_score_))

```

```

Best parameters for this KNN model: {'n_neighbors': 26}
Best score for this KNN model: 0.742

```

```

216 # ROC AUC metrics for explaining the area under the curve
217 |
218 knn_cv.fit(X, y)
219
220 # Compute predicted probabilities: y_pred_prob
221 y_pred_prob = knn_cv.predict_proba(X_test)[:, 1]
222
223 # Compute and print AUC score
224 print("The Area under curve (AUC) on validation dataset is: {:.4f}".format(roc_auc_score(y_test, y_pred_prob)))
225

```

```

The Area under curve (AUC) on validation dataset is: 0.9485

```

```

226 # Compute cross-validated AUC scores: cv_auc
227 cv_auc = cross_val_score(knn_cv, X, y, cv=5, scoring='roc_auc')
228
229 # Print list of AUC scores
230 print("AUC scores computed using 5-fold cross-validation: {}".format(cv_auc))

```

```

AUC scores computed using 5-fold cross-validation: [0.50910538 0.60498652 0.7537216 0.77038827 0.67790078]

```

```

232 # Print CV Scores Mean
233 |
234 print("Cross Validation Scores Mean: {}".format(np.mean(cv_auc)))

```

Cross Validation Scores Mean: 0.6632205108458478

AUC values range from 0 to 1; the AUC for our trained model is 0.9485 or 95%. A model with 100% incorrect predictions has an AUC of 0.0; one with 100% correct predictions has an AUC of 1.0; the greater the AUC, the better the model's performance in distinguishing between positive and negative classes.

## E2. Classification Analysis Results and Implications

The below chart summarizes/compares our intermediate and scaled models. The scaled modeled improved the accuracy of predictions by 11%; the AUC of the scaled model signifies that it performs significantly greater at distinguishing between positive and negative classes.

	Intermediate Model	Scaled Model	-/+
<b>Precision</b>	47%	71%	+ 24%
<b>Recall</b>	38%	58%	+ 20%
<b>Accuracy</b>	71%	82%	+ 11%
<b>CV mean</b>		66%	
<b>AUC</b>		95%	

## E3. Analysis Limitation

Although the scaled model's accuracy increased by 11%, it still does not indicate what commonalities exist between the churn customers; what services they shared; what their customer satisfaction ratings were; and how long they were customers before churning.

## E4. Recommended course of Action

A recommended course of action would be to conduct additional analysis into each of the variables for those churn customers to find commonalities and trends; and it could be a combination of variables (multiple services) that will indicate whether a customer (new or

existing) will churn; this further in-dept analysis will provide management with more hindsight and direction when making decisions on what services to offer, contracts (tenure), and pricing.

**G. Panopto video recording**

[VideoLink](#)

## References

Dwivedi, R. (n.d.). How does K-nearest neighbor works in Machine Learning: KNN algorithm.

How does K-nearest Neighbor Works in Machine Learning | KNN algorithm. Retrieved July 21, 2022, from <https://www.analyticssteps.com/blogs/how-does-k-nearest-neighbor-works-machine-learning-classification-problem>

Google. (n.d.). Classification: Roc curve and AUC | Machine Learning | Google Developers.

Google. Retrieved July 23, 2022, from <https://developers.google.com/machine-learning/crash-course/classification/roc-and-auc>

Nelson, D. (2020, August 23). What is a KNN (K-nearest neighbors)? Unite.AI. Retrieved July

21, 2022, from <https://www.unite.ai/what-is-k-nearest-neighbors/#:~:text=The%20primary%20assumption%20that%20a,two%20points%20on%20a%20graph.>

WGU. (n.d.). NVM2 TASK 1: Classification Analysis. WGU Performance Assessment.

Retrieved July 23, 2022, from

<https://tasks.wgu.edu/student/000194226/course/20900018/task/2807/overview>